


Process State Diagram



Process Creation

- Processes are created when an existing process calls the `fork()` function
- New process created by `fork` is called the *child process*
- `fork()` is a function that is called once but returns twice
 - Only difference in the return value. Returns 0 in the child process and the child's PID in the parent process

fork() man page

FORK(2)

Linux Programmer's Manual

FORK(2)

NAME

fork – create a child process

SYNOPSIS

```
#include <sys/types.h>
#include <unistd.h>
```

```
pid_t fork(void);
```

DESCRIPTION

fork() creates a child process that differs from the parent process only in its PID and PPID, and in the fact that resource utilizations are set to 0. File locks and pending signals are not inherited.

fork() return values

RETURN VALUE

On success, the PID of the child process is returned in the parent's thread of execution, and a 0 is returned in the child's thread of execution. On failure, a -1 will be returned in the parent's context, no child process will be created, and errno will be set appropriately.

ERRORS

EAGAIN `fork()` cannot allocate sufficient memory to copy the parent's page tables and allocate a task structure for the child.

EAGAIN It was not possible to create a new process because the caller's **RLIMIT_NPROC** resource limit was encountered. To exceed this limit, the process must have either the **CAP_SYS_ADMIN** or the **CAP_SYS_RESOURCE** capability.

ENOMEM `fork()` failed to allocate the necessary kernel structures because memory is tight.

Parent and children PIDs

- Why does `fork` return the child's PID to parent processes but it returns 0 to the children?
 - Provides a way to determine which process you are in.
 - Processes can only have one parent. To determine the parent PID you can call `getppid()`.
 - Processes can have multiple children so there is no way for a function to give a parent its child PID

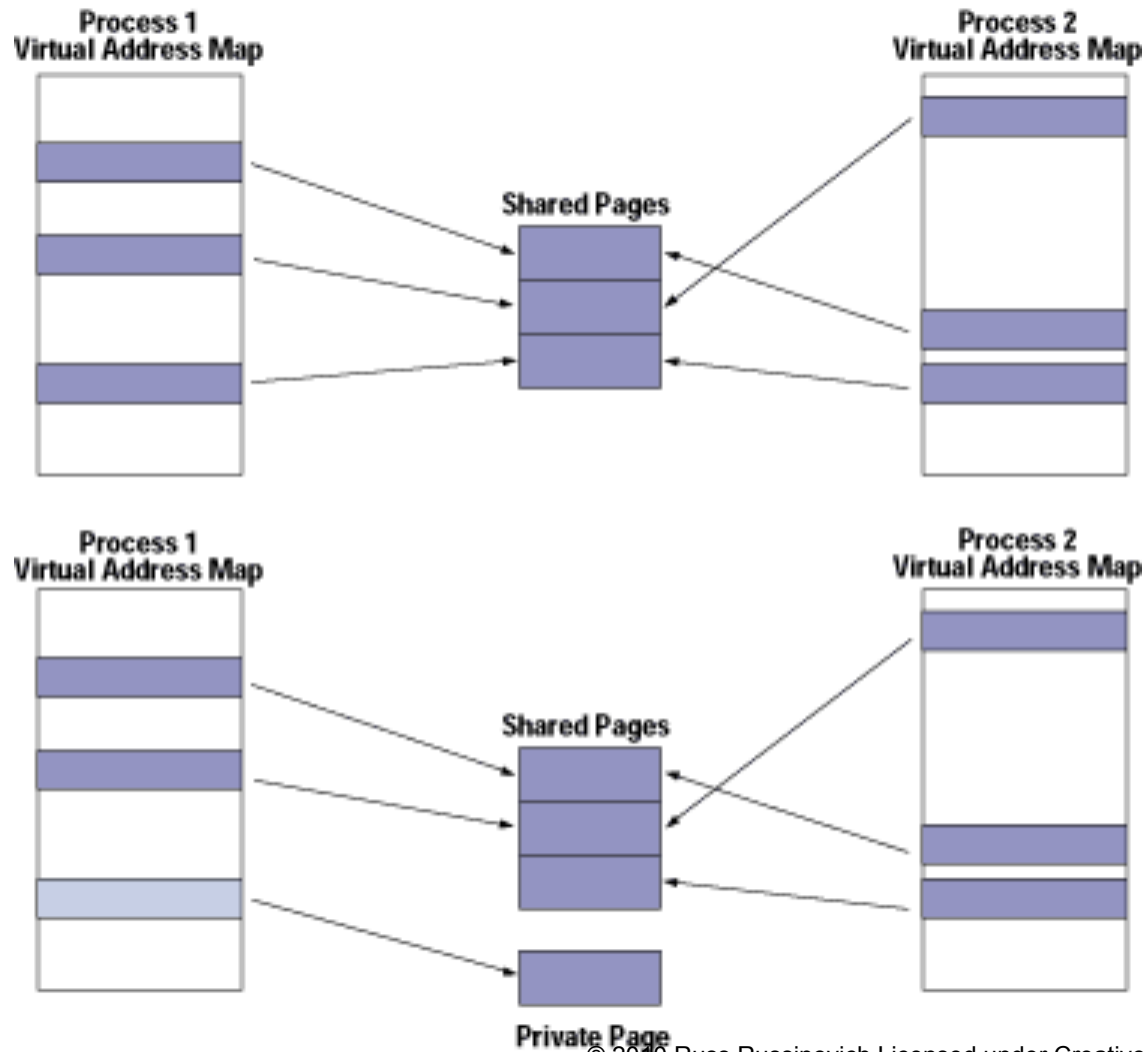
Process Creation

- After `fork()` both the parent and child continue executing with the instruction that follows the call to `fork()`.
- The child process is a copy of the parent process. The child gets:
 - copy of the parent's data space
 - copy of the parent's heap
 - copy of the parent's stack
 - text segment if it's read-only


Copy-On-Write

- On Linux the parent process's pages are not copied for the child process.
- The pages are shared between the child and the parent process.
 - Pages are marked read-only
- When either process modifies a page, a page fault occurs and a separate copy of that particular page is made for that process which performed the modification.
- This process will then use the newly copied page rather than the shared one in all future references. The other process continues to use the original copy of the page

Copy-On-Write



File Control Blocks



Inherited Properties

- user ID, group ID
- process group ID
- controlling terminal
- current working directory
- root directory
- file mode creation mask
- signal mask
- environment
- attached shared memory segments
- resource limits

Unique Properties

- the return value from `fork()`
- the process IDs
- the parent process IDs
- file locks
- pending alarms are cleared for the child
- the set of pending signals for the child is set to zero


fork() code example

```
#include <sys/types.h>
#include <sys/wait.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(void)
{
    pid_t pid = fork();

    if (pid == -1) {
        // When fork() returns -1, an error happened.
        perror("fork failed");
        exit(EXIT_FAILURE);
    }
    else if (pid == 0) {
        // When fork() returns 0, we are in the child process.
        printf("Hello from the child process!\n");
        fflush(NULL);
        exit(EXIT_SUCCESS);
    }
    else {
        // When fork() returns a positive number, we are in the parent process
        // and the return value is the PID of the newly created child process.
        int status;
        (void)waitpid(pid, &status, 0);
        printf("Hello form the parent process!");
        fflush(NULL);
    }
    return EXIT_SUCCESS;
}
```

Hello World



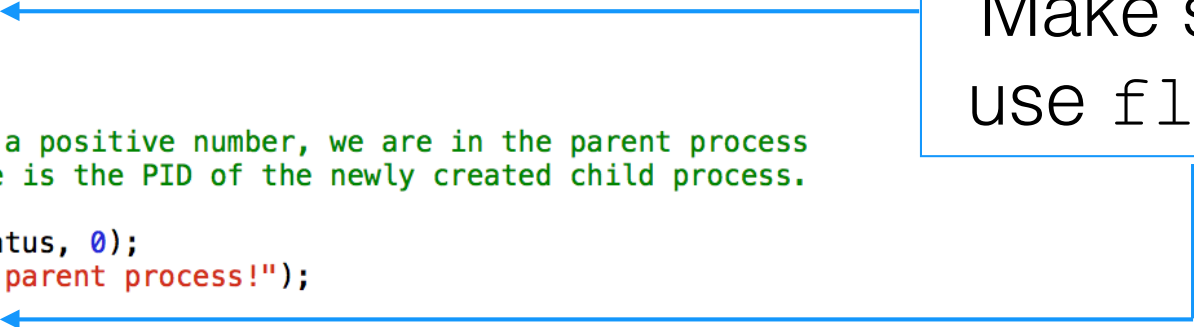
fork() code example

```
#include <sys/types.h>
#include <sys/wait.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(void)
{
    pid_t pid = fork();

    if (pid == -1) {
        // When fork() returns -1, an error happened.
        perror("fork failed");
        exit(EXIT_FAILURE);
    }
    else if (pid == 0) {
        // When fork() returns 0, we are in the child process.
        printf("Hello from the child process!\n");
        fflush(NULL);
        exit(EXIT_SUCCESS);
    }
    else {
        // When fork() returns a positive number, we are in the parent process
        // and the return value is the PID of the newly created child process.
        int status;
        (void)waitpid(pid, &status, 0);
        printf("Hello form the parent process!");
        fflush(NULL);
    }
    return EXIT_SUCCESS;
}
```

I/O is buffered.
Make sure to
use `flush()`

A blue box containing the text "I/O is buffered. Make sure to use flush()". Two blue arrows originate from this box. One arrow points to the `fflush(NULL);` line in the child process block of the code. The other arrow points to the `fflush(NULL);` line in the parent process block of the code.

fork() failures

- Too many processes already.
 - Usually a sign something else has gone wrong
 - `cat /proc/sys/kernel/pid_max`
- Too many processes for the current user
 - `ulimit -a`

Terminating a Process

- Three normal ways
 1. Executing a return from the `main` function
 2. Calling the `exit` function. C library function.
 3. Calling the `_exit` function. System call.
 1. You should use `_exit` to kill the child program when the `exec` fails. Otherwise the child process may interfere with the parent process' external data by calling its signal handlers, or flushing buffers.
 2. You should also use `_exit` in any child process that does not do an `exec`, but those are rare.

Terminating a Process

- Two abnormal ways
 1. Calling `abort`. Generates a SIGABORT signal.
 2. The process received a signal

What if the parent terminates first?

- `exit` and `_exit` return the termination status of child processes to the parent. Who gets the status if the parent has already terminated?
- The `init` process becomes the parent of any process whose parent terminates.
 - The orphaned process is inherited by the `init` process
 - When a process is terminated the kernel iterates through all the active processes to see if the terminating process is the parent of any remaining processes. If so, it inherits that process

Zombie Process

- A process that has completed execution but still has an entry in the process table.
 - The entry is still needed to allow the parent process to read its child's exit status
- A child process always becomes a zombie before being removed from the resource table.
- Normally, zombies are immediately waited on by their parent and then reaped by the system

`wait()` and `waitpid()`

- When a process terminates the parent is notified by the operating system sending a SIGCHLD signal.
 - Default handling is to ignore the signal
- Child termination is asynchronous
- POSIX provides two system calls `wait` and `waitpid`

wait()

WAIT(2)

Linux Programmer's Manual

WAIT(2)

NAME

wait, waitpid – wait for process to change state

SYNOPSIS

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t wait(int *status);
```

```
pid_t waitpid(pid_t pid, int *status, int options);
```

```
int waitid(idtype_t idtype, id_t id, siginfo_t *infop, int options);
```

- `wait` suspends the parent's operation until the next child process terminates.
- If there is a child process already terminated and waiting for reaping, `wait` will return immediately.
- PID of the terminated child process is returned on success. -1 if failure.

wait () status parameter

WIFEXITED(status)

returns true if the child terminated normally, that is, by calling **exit(3)** or **_exit(2)**, or by returning from **main()**.

WEXITSTATUS(status)

returns the exit status of the child. This consists of the least significant 16-8 bits of the status argument that the child specified in a call to **exit()** or **_exit()** or as the argument for a return statement in **main()**. This macro should only be employed if **WIFEXITED** returned true.

WIFSIGNALED(status)

returns true if the child process was terminated by a signal.

WTERMSIG(status)

returns the number of the signal that caused the child process to terminate. This macro should only be employed if **WIFSIGNALED** returned true.

wait () status parameter

WCOREDUMP(status)

returns true if the child produced a core dump. This macro should only be employed if **WIFSIGNALED** returned true. This macro is not specified in POSIX.1-2001 and is not available on some Unix implementations (e.g., AIX, SunOS). Only use this enclosed in `#ifdef WCOREDUMP ... #endif`.

WIFSTOPPED(status)

returns true if the child process was stopped by delivery of a signal; this is only possible if the call was done using **WUNTRACED** or when the child is being traced (see **ptrace(2)**).

WSTOPSIG(status)

returns the number of the signal which caused the child to stop. This macro should only be employed if **WIFSTOPPED** returned true.

WIFCONTINUED(status)

(Since Linux 2.6.10) returns true if the child process was resumed by delivery of **SIGCONT**.

waitpid()

```
pid_t waitpid(pid_t pid, int *status, int options);
```

- Allows the parent process to wait on a specific child process.

wait () code example

```
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
#include <stdio.h>
#include <string.h>

int main(void)
{
    pid_t child_pid = fork();
    int status;

    if (child_pid == 0)
    {
        // Sleep for a second
        sleep(1);

        // Intentionally SEGFAULT the child process
        int *p = NULL;
        *p = 1;

        exit(0);
    }

    // Wait for the child to exit
    waitpid( child_pid, &status, 0 );

    // See if the child was terminated by a signal
    if( WIFSIGNALED( status ) )
    {
        // Print the signal that the child terminated with
        printf("Child returned with status %d\n", WTERMSIG( status ) );
    }

    return 0;
}
```

Process Creation

`fork()` creates an exact copy of a process. How do we create a unique process?

By combining `fork` with an `exec` function

exec functions

- When exec is called the new program, specified by exec, completely replaces the running process.
 - text, data, heap and stack are all replaced
 - PID stays the same since it's not a new process

exec family

```
int execl(char const *path, char const *arg0, ...);  
int execlp(char const *path, char const *arg0, ..., char const *envp[]);  
int execlp(char const *file, char const *arg0, ...);  
int execv(char const *path, char const *argv[]);  
int execve(char const *path, char const *argv[], char const *envp[]);  
int execvp(char const *file, char const *argv[]);
```

Meaning of the letters after exec:

l – A list of command-line arguments are passed to the function.

e – An array of pointers to environment variables is passed to the new process image.

p – The PATH environment variable is used to find the file named in the path argument.

v – Command-line arguments are passed to the function as an array of pointers.

exec code example


```
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
#include <stdio.h>

int main(void)
{
    pid_t child_pid = fork();
    int status;

    if (child_pid == 0)
    {
        execl("/bin/ls", "ls", NULL );
        exit(0);
    }


    // Wait for the child to exit
    waitpid( child_pid, &status, 0 );

    return 0;
}
```



Signals

- A signal is a asynchronous notification of an event
- Signals are how the operating system communicates with an application process.
- When a signal is received the process stops running and the assigned signal handler is run.
- When the signal has been handled the process resumes where it left off.



Some examples

- SIGSEGV - Known as a segmentation fault. It occurs when a process makes an illegal memory reference. The process halts and the default signal handler exits the process.
- SIGINT - The interrupt signal occurs when the user types ctrl-c. The default signal handler will exit the process.

Sending Signals

- By keyboard
 - ctrl-c sends the SIGINT signal and the receiving process exits
 - ctrl-z sends the SIGTSTP signal and the receiving process is suspended
 - ctrl-\ sends the SIGQUIT signal and the receiving process exits

Sending signals via shell commands

`kill -signal pid`

- Send a signal, specified by `-signal`, to the process with the process id `pid`.
- If no signal is specified then the signal sent is `SIGTERM`

Example:

`kill -9 3412`

`kill -SIGQUIT 3412`

Send signal via a function

```
int kill( pid_t pid, int sig)
```

- Send a signal, specified by -signal, to the process with the process id pid.
- If no signal is specified then the signal sent is SIGTERM

Example:

```
pid_t pid = getpid();    // process gets its own pid  
kill(pid, SIGINT);       // and sends itself a SIGINT
```

Send signal via a function

```
int raise( int sig )
```

- Send a signal to the current process

Example:

```
raise(SIGINT); // process sends itself a SIGINT
```

Signal Numbers

```
[bakker@crystal ~]$ kill -l
```

1) SIGHUP	2) SIGINT	3) SIGQUIT	4) SIGILL
5) SIGTRAP	6) SIGABRT	7) SIGBUS	8) SIGFPE
9) SIGKILL	10) SIGUSR1	11) SIGSEGV	12) SIGUSR2
13) SIGPIPE	14) SIGALRM	15) SIGTERM	16) SIGSTKFLT
17) SIGCHLD	18) SIGCONT	19) SIGSTOP	20) SIGTSTP
21) SIGTTIN	22) SIGTTOU	23) SIGURG	24) SIGXCPU
25) SIGXFSZ	26) SIGVTALRM	27) SIGPROF	28) SIGWINCH
29) SIGIO	30) SIGPWR	31) SIGSYS	34) SIGRTMIN
35) SIGRTMIN+1	36) SIGRTMIN+2	37) SIGRTMIN+3	38) SIGRTMIN+4
39) SIGRTMIN+5	40) SIGRTMIN+6	41) SIGRTMIN+7	42) SIGRTMIN+8
43) SIGRTMIN+9	44) SIGRTMIN+10	45) SIGRTMIN+11	46) SIGRTMIN+12
47) SIGRTMIN+13	48) SIGRTMIN+14	49) SIGRTMIN+15	50) SIGRTMAX-14
51) SIGRTMAX-13	52) SIGRTMAX-12	53) SIGRTMAX-11	54) SIGRTMAX-10
55) SIGRTMAX-9	56) SIGRTMAX-8	57) SIGRTMAX-7	58) SIGRTMAX-6
59) SIGRTMAX-5	60) SIGRTMAX-4	61) SIGRTMAX-3	62) SIGRTMAX-2
63) SIGRTMAX-1	64) SIGRTMAX		

Signal Terminology

- A signal is *generated* for a process when the event that causes the signal occurs. The event can be a hardware exception, a software condition, or a call to the kill function
- When a signal is generated the kernel usually sets a flag in the process table

Signal Terminology

- A signal is *delivered* to a process when the action for a signal is taken.
- During the time between the generation of the signal and the delivery of the signal the signal is said to be *pending*.

Signal Handling

- Every signal is assigned a default handler. Usually they just exit the process.
- Programs can install their own signal handlers for most signals.
- Can't install handlers for:
 - SIGKILL
 - SIGSTOP

Installing a signal handler

SIGACTION(2)

Linux Programmer's Manual

SIGACTION(2)

NAME

`sigaction` – examine and change a signal action

SYNOPSIS

```
#include <signal.h>
```

```
int sigaction(int signum, const struct sigaction *act, struct sigaction *oldact);
```

- The function *handler* specified in the `sigaction` struct is installed as the new handler for the signal specified by `signum`.
- The new handler is called every time the process receives a signal of type `signum`.

Signal Handler Example

- See `sigint.c` and `multiple_signal_handlers.c` on course website under Code Samples

Blocking Signals

- We can block signals, but why would we want to?
 - Race conditions. What happens when a signal is received while we are in the middle of handling the same type of signal?
- The POSIX standard provides a way to block signals.

Blocking Signals

- Each process has a signal mask
 - The operating system uses the mask to determine which signals to deliver.
- The function `sigprocmask()` provides a way for programs to modify their signal mask to block and unblock signals.

Blocking Signals

- If a signal is generated for a process but blocked then the signal remains pending for the process until the process either:
 - unblocks the signal
 - changes the action to ignore the signal

Multiple Blocked Signals

- The POSIX standard allows the OS to deliver one or multiple copies of the signal.
- Most UNIX implementations do not queue signals and just deliver a single one.

sigprocmask()

SIGPROCMASK(2)

BSD System Calls Manual

SIGPROCMASK(2)

NAME

sigprocmask -- manipulate current signal mask

SYNOPSIS

#include <signal.h>

int
sigprocmask(int how, const sigset_t *restrict set, sigset_t *restrict oset);

- First parameter is *how*. It can be:
 - SIG_BLOCK - adds the provided set to the current signal mask
 - SIG_UNBLOCK - removes the provided set from the current signal mask
 - SIG_SETMASK - the current mask is replaced by the provided set

sigprocmask()

SIGPROCMASK(2)

BSD System Calls Manual

SIGPROCMASK(2)

NAME

sigprocmask -- manipulate current signal mask

SYNOPSIS

#include <signal.h>

int
sigprocmask(int *how*, const sigset_t *restrict *set*, sigset_t *restrict *oset*);

- Second parameter is a pointer to the a signal mask called *set*.
- Third parameter is a pointer to which the old set can be returned.

Signal Blocking Example

- See `sig_set_example.c` on course website under Code Samples

Process Control

